



Emergence of belief-state copies during training

Layer 0 Layer 1 Layer 2 Layer 3


Step 0 (random init)




Step 100



Step 2000



Best model (~46k steps)



Probe iteration

Probe iteration